

CS202 Design Pattern Report: Interpreter Pattern

Formal Grammars, Abstract Syntax Trees, and Modern C++ Memory
Architectures

Group 52

Nguyễn Đình Minh Huy (Student ID: 25125083)

Trần Gia Huy (Student ID: 25125084)

August 27, 2026

Contents

1	Real-world Problem & Motivation	3
1.1	Background: Domain-Specific Languages and Boolean Logic	3
1.2	Problem Statement	3
1.3	Language Grammar and Operator Rules	4
2	Naive Solution	5
2.1	Step-by-Step Execution of the Naive Evaluator	6
3	Analysis of Naive Solution Drawbacks	7
4	Introduction to the Interpreter Design Pattern	7
4.1	Definition and Intent	7
4.2	Key Structural Components	7
5	UML Class Diagrams	9
5.1	General GoF Interpreter Pattern Structure	9
5.2	Boolean Logic Evaluator Class Hierarchy	9
6	Pattern-Based Implementation in C++	10
6.1	Architectural Design Rationale: Strict AST vs. ASG/DAG	10
6.2	Source Code Implementation	10
6.3	Execution Walkthrough & Step-by-Step AST Evaluation Trace	14
7	Evaluation of Trade-offs	15
7.1	Advantages	15
7.2	Disadvantages and Limitations	15
8	Applications	16

9 Observer and Mediator Integration	17
9.1 Combining Interpreter with the Observer Pattern	17
9.2 Combining Interpreter with the Mediator Pattern	17
10 Interactive Quiz	17
11 Conclusion	19
References	20

1 Real-world Problem & Motivation

1.1 Background: Domain-Specific Languages and Boolean Logic

In software engineering, systems frequently encounter requirements to interpret, evaluate, or execute statements written in domain-specific mini-languages at runtime. Unlike general-purpose programming languages (e.g., C++, Java, Rust), domain-specific languages (DSLs) are tailored to specialized problem domains. Prominent examples include:

- **Dynamic Business Rule Engines:** Validating loan eligibility, risk scoring, or discount tiers defined by non-developer business analysts.
- **Database Query and Filter Parsers:** Evaluating SQL-like WHERE clauses (e.g., `age > 21 AND status == 'ACTIVE'`) dynamically generated from user search interfaces.
- **IoT Sensor Event Processing:** Filtering high-throughput telemetry streams against custom user-defined logic conditions (e.g., `temp > 85.0 OR (pressure > 120.0 AND NOT valve_open)`).
- **Configuration-Driven Access Control (RBAC/ABAC):** Evaluating permission predicates such as `(isOwner OR isAdmin) AND NOT isBanned`.

Hardcoding every permutation of logical rules into static source code results in severe inflexibility. Every rule change requires recompilation, regression testing, and deployment. Consequently, systems require a mechanism to define, compose, and evaluate arbitrary sentences of a well-defined language dynamically at runtime.

1.2 Problem Statement

To illustrate the problem concretely, consider a configurable logic engine that evaluates Boolean expressions consisting of:

1. **Variables:** Identifiers representing dynamic state (e.g., A, B, C).
2. **Constants:** Literal truth values (`true`, `false`).
3. **Unary Operators:** Logical negation (`NOT` / $\neg$).
4. **Binary Operators:** Conjunction (`AND` / $\wedge$), Disjunction (`OR` / $\vee$), and Exclusive-OR (`XOR` / $\oplus$).
5. **Groupings:** Nested parentheses controlling operator precedence.

The engine must take an expression along with a symbol table (a *Context* mapping identifiers to their Boolean values) and compute the exact truth result according to standard mathematical logic.

1.3 Language Grammar and Operator Rules

To keep the design clear and practical, the language is defined by simple composition rules rather than formal compiler notation:

- **Atomic Expressions (Terminals):** Either a named variable (e.g., A, B) looked up from the `Context` or a literal Boolean constant (`true`, `false`).
- **Unary Operation:** Negation (`NOT`) operating on a single child expression.
- **Binary Operations:** Conjunction (`AND`), Disjunction (`OR`), and Exclusive-OR (`XOR`) combining two child sub-expressions.
- **Operator Precedence:** Evaluated in standard algebraic order: `NOT` > `AND` > `XOR` > `OR`, with nested sub-expressions enclosed in parentheses evaluated first.

By organizing expressions according to these simple rules, any valid logical statement can be translated directly into a hierarchical tree structure.

2 Naive Solution

In a naive C++ implementation, expressions are treated as raw text strings. The evaluation is performed through manual string searching, linear token splitting, or ad-hoc procedural loops.

```
1 // Context: Maps variable names to boolean values
2 class Context {
3 public:
4     void assign(const std::string& name, bool value) { variables_[name]
      = value; }
5     bool lookup(const std::string& name) const {
6         auto it = variables_.find(name);
7         if (it == variables_.end()) throw std::runtime_error("Unknown
      variable: " + name);
8         return it->second;
9     }
10 private:
11     std::map<std::string, bool> variables_;
12 };
13
14 class NaiveLogicEvaluator {
15 public:
16     // Evaluates flat left-to-right expressions (e.g. "A AND B OR C")
17     static bool evaluate(const std::string& expression, const Context&
      context) {
18         std::vector<std::string> tokens = tokenize(expression);
19         if (tokens.empty()) return false;
20
21         bool result = parseOperand(tokens[0], context);
22
23         for (size_t i = 1; i + 1 < tokens.size(); i += 2) {
24             const std::string& op = tokens[i];
25             bool rhs = parseOperand(tokens[i + 1], context);
26
27             if (op == "AND") {
28                 result = result && rhs;
29             } else if (op == "OR") {
30                 result = result || rhs;
31             } else if (op == "XOR") {
32                 result = result ^ rhs;
33             } else {
34                 throw std::invalid_argument("Unsupported operator: " +
      op);
35             }
36         }
37         return result;
38     }
39 private:
40     static bool parseOperand(const std::string& token, const Context&
      context) {
41         if (token == "true" || token == "1") return true;
42         if (token == "false" || token == "0") return false;
43         return context.lookup(token);
44     }
45 }
46
```

```

47     static std::vector<std::string> tokenize(const std::string&
      expression) {
48         std::stringstream ss(expression);
49         std::string token;
50         std::vector<std::string> tokens;
51         while (ss >> token) tokens.push_back(token);
52         return tokens;
53     }
54 };

```

Listing 1: Naive Boolean Expression Evaluator (`naive.cpp`)

2.1 Step-by-Step Execution of the Naive Evaluator

To illustrate how the naive solution operates, consider evaluating the string "A OR B AND C" with context $A = \text{true}$, $B = \text{false}$, $C = \text{true}$:

1. **Tokenization (String Splitting):** The evaluator parses the raw string by whitespace into a linear token vector:

$$\text{tokens} = ["A", "OR", "B", "AND", "C"]$$

2. **Initial Left Operand Lookup:** The first token "A" is looked up in Context ($A = \text{true}$) to initialize the running accumulator `result = true`.
3. **Linear Left-to-Right Evaluation Loop:** The loop steps through consecutive operator-operand pairs:

- *Step 3a:* Reads operator "OR" and operand "B" ($B = \text{false}$). Updates:

$$\text{result} = \text{result} \vee B = \text{true} \vee \text{false} = \text{true}$$

- *Step 3b:* Reads operator "AND" and operand "C" ($C = \text{true}$). Updates:

$$\text{result} = \text{result} \wedge C = \text{true} \wedge \text{true} = \text{true}$$

Critical Flaw Demonstrated: The naive evaluator computes $(A \vee B) \wedge C = (\text{true} \vee \text{false}) \wedge \text{true} = \text{true}$. If the input were "false OR true AND false", the mathematically correct precedence $A \vee (B \wedge C)$ yields $\text{false} \vee (\text{true} \wedge \text{false}) = \text{false}$, but the naive left-to-right evaluator incorrectly outputs $(\text{false} \vee \text{true}) \wedge \text{false} = \text{false}$ (or in other expressions, completely corrupts truth values due to lack of precedence and parenthesis support).

3 Analysis of Naive Solution Drawbacks

The naive string-scanning and procedural dispatch approach presents several architectural, correctness, and performance limitations:

Deficiency Area	Manifestation in Naive Code	Architectural Consequence
Open-Closed Principle (OCP) Violation	Adding an operator (e.g., NAND, IMPLIES) requires modifying the central <code>if-else</code> loop in <code>evaluate()</code> .	Existing evaluation code must be modified, increasing regression risk.
Single Responsibility Principle (SRP) Violation	<code>NaiveLogicEvaluator</code> handles lexical tokenization, semantic validation, operator precedence, and evaluation in one class.	Parsing and evaluation logic cannot be tested or maintained independently.
Operator Precedence Limitations	Evaluates flat expressions strictly left-to-right. <code>A OR B AND C</code> evaluates as $(A \vee B) \wedge C$ instead of $A \vee (B \wedge C)$.	Produces incorrect results for expressions requiring operator precedence.
Nesting and Parentheses Support	Supporting nested sub-expressions (e.g., <code>(A AND (NOT B)) OR C</code>) requires recursive regex or complex stack tracking.	Significantly increases procedural string manipulation complexity.
Evaluation Overhead	String allocations, repeated <code>substr()</code> calls, and linear scans incur additional overhead on repeated evaluations.	Inefficient when expressions are evaluated repeatedly against changing contexts.

Table 1: Design and Performance Limitations of the Naive Solution

4 Introduction to the Interpreter Design Pattern

4.1 Definition and Intent

According to the Gang of Four (GoF), the **Interpreter Pattern** is defined as:

"Given a language, define a representation for its grammar along with an interpreter that uses the representation to interpret sentences in the language."

The pattern addresses the problem by representing each grammar rule as a class. The syntax of an expression is represented as an **Abstract Syntax Tree (AST)**, where each node is an instance of an expression class. Evaluation is performed by traversing the tree recursively.

4.2 Key Structural Components

- **Context (Context):** Encapsulates the global execution environment, maintaining state such as variable bindings (symbol table) needed by terminal expressions.

- **AbstractExpression** (BooleanExpression): Declares an abstract interface with an evaluation method:

```
virtual bool interpret(const Context& context) const = 0;
```

- **TerminalExpression** (VariableExpression, ConstantExpression): Represents leaf nodes corresponding to terminal grammar symbols. Directly evaluates based on the provided Context.
- **NonterminalExpression** (AndExpression, OrExpression, NotExpression, XorExpression): Represents internal nodes of the AST corresponding to grammar production rules. Contains child **AbstractExpression** instances and aggregates their recursive evaluations.
- **Client**: Assembles (or receives from a parser) the composite AST hierarchy and initiates evaluation by invoking `interpret()` on the root node.

5 UML Class Diagrams

5.1 General GoF Interpreter Pattern Structure

The Interpreter pattern defines a polymorphic hierarchy where non-terminal composite nodes hold references to abstract expressions, while terminal nodes evaluate directly against the context.




Figure 1: General GoF Interpreter Pattern UML Structure

5.2 Boolean Logic Evaluator Class Hierarchy

For our Boolean expression engine, **BooleanExpression** acts as the abstract component interface with concrete terminals (**VariableExpression**, **ConstantExpression**) and non-terminals owning child expressions via `std::unique_ptr`.




Figure 2: Concrete Boolean Logic Evaluator Class Hierarchy

6 Pattern-Based Implementation in C++17

6.1 Architectural Design Rationale: Strict AST vs. ASG/DAG

A key design decision in our C++ implementation is enforcing a strict **Abstract Syntax Tree (AST)** rather than an **Abstract Semantic Graph (ASG / DAG)**:

- **Strict Tree Structure** (`std::unique_ptr`): Every expression node has exactly *one unique parent*. Duplicate occurrences of a variable (such as A appearing in multiple sub-clauses) are instantiated as distinct terminal leaves. This is enforced using `std::unique_ptr` with move semantics, eliminating reference counting overhead and ensuring deterministic single-owner memory reclamation.
- **Graph Alternative** (`std::shared_ptr`): In optimizing engines with Common Subexpression Elimination (CSE), identical sub-expressions share node instances with multiple parents (a DAG). However, this introduces atomic reference counting overhead and cache complexity.

6.2 Source Code Implementation

The following listings demonstrate the complete, robust implementation using C++17 move semantics and unique pointers (`Interpreter/src/pattern.cpp`).

```

1 #include <iostream>
2 #include <string>
3 #include <map>
4 #include <memory>
5 #include <stdexcept>
6
7 namespace pattern {
8
9 // Context: Holds runtime symbol bindings
10 class Context {
11 public:
12     void assign(const std::string& name, bool value) {

```

```

13     variables_[name] = value;
14 }
15
16 bool lookup(const std::string& name) const {
17     auto it = variables_.find(name);
18     if (it == variables_.end()) {
19         throw std::runtime_error("Unbound variable in Context: " +
20     name);
21     }
22     return it->second;
23 }
24 private:
25     std::map<std::string, bool> variables_;
26 };
27
28 // Abstract Expression Interface
29 class BooleanExpression {
30 public:
31     virtual ~BooleanExpression() = default;
32     virtual bool interpret(const Context& context) const = 0;
33 };

```

Listing 2: Context and Abstract Expression Interface

```

1 // Terminal Expression: Variable Lookup
2 class VariableExpression : public BooleanExpression {
3 public:
4     explicit VariableExpression(std::string name) : name_(std::move(
5     name)) {}
6
7     bool interpret(const Context& context) const override {
8         return context.lookup(name_);
9     }
10 private:
11     std::string name_;
12 };
13
14 // Terminal Expression: Boolean Literal Constant
15 class ConstantExpression : public BooleanExpression {
16 public:
17     explicit ConstantExpression(bool value) : value_(value) {}
18
19     bool interpret(const Context&) const override {
20         return value_;
21     }
22 private:
23     bool value_;
24 };

```

Listing 3: Terminal Expressions: Variables and Constants

```

1 // Non-Terminal: Logical Conjunction (AND)
2 class AndExpression : public BooleanExpression {
3 public:
4     AndExpression(std::unique_ptr<BooleanExpression> left, std::
unique_ptr<BooleanExpression> right)
5         : left_(std::move(left)), right_(std::move(right)) {}
6
7     bool interpret(const Context& context) const override {
8         return left_->interpret(context) && right_->interpret(context);
9     }
10
11 private:
12     std::unique_ptr<BooleanExpression> left_;
13     std::unique_ptr<BooleanExpression> right_;
14 };
15
16 // Non-Terminal: Logical Disjunction (OR)
17 class OrExpression : public BooleanExpression {
18 public:
19     OrExpression(std::unique_ptr<BooleanExpression> left, std::
unique_ptr<BooleanExpression> right)
20         : left_(std::move(left)), right_(std::move(right)) {}
21
22     bool interpret(const Context& context) const override {
23         return left_->interpret(context) || right_->interpret(context);
24     }
25
26 private:
27     std::unique_ptr<BooleanExpression> left_;
28     std::unique_ptr<BooleanExpression> right_;
29 };
30
31 // Non-Terminal: Logical Negation (NOT)
32 class NotExpression : public BooleanExpression {
33 public:
34     explicit NotExpression(std::unique_ptr<BooleanExpression> operand)
35         : operand_(std::move(operand)) {}
36
37     bool interpret(const Context& context) const override {
38         return !operand_->interpret(context);
39     }
40
41 private:
42     std::unique_ptr<BooleanExpression> operand_;
43 };
44
45 // Non-Terminal: Exclusive-OR (XOR)
46 class XorExpression : public BooleanExpression {
47 public:
48     XorExpression(std::unique_ptr<BooleanExpression> left, std::
unique_ptr<BooleanExpression> right)
49         : left_(std::move(left)), right_(std::move(right)) {}
50
51     bool interpret(const Context& context) const override {
52         return left_->interpret(context) ^ right_->interpret(context);
53     }
54
55 private:

```

```
56     std::unique_ptr<BooleanExpression> left_;  
57     std::unique_ptr<BooleanExpression> right_;  
58 };  
59  
60 } // namespace pattern
```

Listing 4: Non-Terminal Expressions: AND, OR, NOT, XOR with Unique Ownership

6.3 Execution Walkthrough & Step-by-Step AST Evaluation Trace

Consider the composite logical expression:

$$\text{Expression} = (A \wedge (\neg B)) \vee (C \oplus A) \quad (1)$$

Given the runtime variable bindings in Context:

$$A = \text{true}, \quad B = \text{false}, \quad C = \text{true}$$

```
1 void runPatternDemo() {
2     pattern::Context context;
3     context.assign("A", true);
4     context.assign("B", false);
5     context.assign("C", true);
6
7     // Build AST: (A AND (NOT B)) OR (C XOR A)
8     auto notB = std::make_unique<pattern::NotExpression>(
9         std::make_unique<pattern::VariableExpression>("B")
10    );
11    auto aAndNotB = std::make_unique<pattern::AndExpression>(
12        std::make_unique<pattern::VariableExpression>("A"),
13        std::move(notB)
14    );
15    auto cXorA = std::make_unique<pattern::XorExpression>(
16        std::make_unique<pattern::VariableExpression>("C"),
17        std::make_unique<pattern::VariableExpression>("A")
18    );
19    auto fullExpr = std::make_unique<pattern::OrExpression>(
20        std::move(aAndNotB),
21        std::move(cXorA)
22    );
23
24    bool result = fullExpr->interpret(context);
25    std::cout << "Result: " << (result ? "true" : "false") << "\n"; //
26    Outputs: true
27 }
```

Listing 5: Building and Evaluating the AST

Table 2 summarizes the post-order depth-first traversal and evaluation sequence of the AST nodes:

Step	Node Evaluated	Subtree	Value	Explanation / Operation
1	VariableExpression("B")	B	false	Context lookup returns $B = \text{false}$.
2	NotExpression	$\neg B$	true	Negation: $!\text{false} = \text{true}$.
3	VariableExpression("A")	A	true	Context lookup returns $A = \text{true}$.
4	AndExpression	$A \wedge (\neg B)$	true	Conjunction: $\text{true} \wedge \text{true} = \text{true}$.
5	VariableExpression("C")	C	true	Context lookup returns $C = \text{true}$.
6	VariableExpression("A")	A	true	Context lookup returns $A = \text{true}$.
7	XorExpression	$C \oplus A$	false	Exclusive OR: $\text{true} \oplus \text{true} = \text{false}$.
8	OrExpression (Root)	$(A \wedge \neg B) \vee (C \oplus A)$	true	Root Disjunction: $\text{true} \vee \text{false} = \text{true}$.

Table 2: Recursive Post-Order Traversal Trace for AST Evaluation

7 Evaluation of Trade-offs

7.1 Advantages

- **High Grammar Extensibility (OCP Compliance):** Introducing new language constructs or operators (e.g., adding an `ImpliesExpression` for $P \implies Q \equiv \neg P \vee Q$) requires creating a new subclass of `BooleanExpression` without altering any existing classes.
- **Single Responsibility Principle (SRP):** Each class is strictly responsible for interpreting exactly one grammar symbol.
- **Decoupled Syntax Structure from Execution:** The structural representation (AST) is cleanly separated from the runtime evaluation state (`Context`).
- **Zero Evaluation Overhead with `std::unique_ptr`:** By relying on direct polymorphic virtual dispatch and unique pointers, tree evaluation avoids string-slicing allocations and runs at native CPU speed.

7.2 Disadvantages and Limitations

- **Combinatorial Class Explosion for Large Grammars:** For complex programming languages with hundreds of grammar production rules (e.g., full ECMAScript or SQL-92), maintaining hundreds of distinct AST classes manually becomes unmanageable.
- **Call-Stack Depth and Recursion Overhead:** Deeply nested syntax trees can lead to stack overflow if evaluated via naive recursion on resource-constrained embedded systems.

- **Grammar Maintenance Overhead:** Modifying operator precedence rules requires structural modifications to the AST assembly rather than simple configuration changes.

8 Applications

The Interpreter pattern is widely applied in software engineering when a system needs to parse, evaluate, or execute domain-specific grammar rules:

- **Domain-Specific Query and Filter Languages:** Database abstraction layers and search systems parse user-defined filter expressions (such as SQL `WHERE` clauses or custom search query languages) into syntax trees to evaluate records dynamically.
- **Business Rule and Validation Engines:** Enterprise workflow platforms use domain-specific mini-languages to evaluate eligibility criteria, discount pricing tiers, or risk scores defined by non-programmer domain experts at runtime.
- **Spreadsheet Formula Evaluation:** Spreadsheet engines parse formula strings (such as mathematical formulas, logical branching, and cell aggregations) into Abstract Syntax Trees that re-evaluate when referenced cell values change.
- **Template Processing and Formatting Engines:** Document and string template engines interpret dynamic placeholder expressions, conditional blocks, and iterative tags embedded in text or markup.
- **Command-Line and Script Interpreters:** Interactive console utilities and automation engines parse script commands, parameterized flags, and compound logical pipelines into executable syntax structures.

9 Observer and Mediator Integration

In real-world software architectures, design patterns rarely exist in isolation. The Interpreter pattern naturally combines with the Observer and Mediator patterns to create reactive, event-driven, and highly decoupled systems.

9.1 Combining Interpreter with the Observer Pattern

In a static Interpreter implementation, clients must manually invoke `interpret()` whenever variable values change. In modern reactive architectures—such as reactive UI frameworks, spreadsheet dependency graphs, or IoT sensor triggers—this workflow is refactored by applying the Observer pattern.

The evaluation `Context` functions as an observable Subject. When a variable's value is modified (for example, via `Context::assign("temperature", 92.5)`), the Context publishes a state-change notification. Root expressions and dependent components register as Observers listening for specific variable keys. Upon receiving an update event, they automatically trigger re-interpretation of their AST and emit new evaluation results reactively, eliminating wasteful CPU polling loops.

9.2 Combining Interpreter with the Mediator Pattern

In complex distributed systems, such as air traffic control systems, command dispatchers, or workflow engines, components must coordinate without direct peer-to-peer coupling. Here, the Mediator pattern leverages an embedded Interpreter engine.

Colleague components do not interact directly with one another. Instead, they send declarative query or command expressions to a central `Mediator`. The central Mediator uses an internal Interpreter engine to dynamically parse incoming expressions, evaluate business rules, and check permission predicates before routing actions to target colleagues. This decouples the communication topology into an $O(N)$ star architecture while retaining the flexibility to modify business coordination rules dynamically at runtime.

10 Interactive Quiz

1. What is the primary intent of the Interpreter Design Pattern?
 - A) To define a family of algorithms, encapsulate each one, and make them interchangeable.
 - B) Given a language, define a representation for its grammar along with an interpreter that uses the representation to interpret sentences in the language.
 - C) To attach additional responsibilities to an object dynamically without subclassing.
 - D) To ensure a class has only one instance and provide a global access point.
2. What is the fundamental difference between a `TerminalExpression` and a `NonTerminalExpression`?
 - A) Terminal expressions represent leaf nodes in an AST (e.g., variables, constants) and evaluate directly; `NonTerminal` expressions represent grammar operators (e.g., AND, OR) containing child sub-expressions.

- B) Terminal expressions only run in console terminals.
 - C) NonTerminal expressions cannot access the `Context` symbol table.
 - D) Terminal expressions require multi-threading while NonTerminal expressions are single-threaded.
3. What is the primary responsibility of the `Context` object in an Interpreter implementation?
- A) To render the GUI window for the user.
 - B) To store global state, symbol table mappings (e.g., variable identifiers to boolean values), and evaluation environment information outside the syntax tree.
 - C) To compile the C++ source code into binary machine code.
 - D) To allocate memory using raw `malloc` calls.
4. In modern C++17, why is `std::unique_ptr` the preferred smart pointer for holding child expressions in non-terminal AST nodes?
- A) Because AST trees exhibit strict single-ownership hierarchy with no sharing between distinct parent nodes, guaranteeing automatic deep deletion with zero reference counting overhead.
 - B) Because `std::shared_ptr` cannot hold abstract polymorphic types.
 - C) Because `std::unique_ptr` automatically compiles expressions to assembly.
 - D) Because `std::weak_ptr` is deprecated in C++17.
5. Under which scenario is the Interpreter pattern generally NOT recommended?
- A) When evaluating simple domain-specific filter queries.
 - B) When implementing dynamic business rule engines.
 - C) When parsing complex, large-scale programming language grammars with hundreds of production rules (where Parser Generators like ANTLR or Bison are more maintainable).
 - D) When parsing boolean logic formulas.
6. How does integrating the Observer pattern enhance an Interpreter AST evaluation pipeline?
- A) It eliminates the need for virtual destructor definitions.
 - B) The `Context` acts as an observable Subject that automatically triggers re-evaluation of dependent AST root expressions whenever variable values change.
 - C) It converts the AST tree into a linear linked list.
 - D) It enforces alphabetical ordering on terminal variable lookups.

Answer Key

1-B, 2-A, 3-B, 4-A, 5-C, 6-B.

11 Conclusion

The Interpreter Design Pattern provides an elegant, object-oriented solution for representing and evaluating dynamic domain-specific grammars. By modeling production rules as distinct classes and organizing them into an Abstract Syntax Tree (AST), the pattern achieves strict adherence to the Open-Closed and Single Responsibility Principles.

Furthermore, combining the Interpreter pattern with modern C++17 memory management idioms—specifically using `std::unique_ptr` for exclusive node ownership and zero-overhead move semantics—yields a high-performance, memory-safe execution engine. When combined with the Observer and Mediator patterns, it forms the backbone for reactive rule engines and decoupled, dynamic event-driven architectures.

References

- Gamma, E., Helm, R., Johnson, R., & Vlissides, J. (1994). *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley.
- Aho, A. V., Lam, M. S., Sethi, R., & Ullman, J. D. (2006). *Compilers: Principles, Techniques, and Tools* (2nd ed.). Addison-Wesley.
- Meyers, S. (2014). *Effective Modern C++: 42 Specific Ways to Improve Your Use of C++11 and C++14*. O'Reilly Media.
- Refactoring.Guru. *Design Patterns: Interpreter Pattern*. <https://refactoring.guru/design-patterns/interpreter>